AdaBoost Does Not Always Cycle

March 31, 2026

Abstract

We give a computer-assisted counterexample to the open question, posed by Rudin, Schapire, and Daubechies in COLT 2012, of whether exhaustive AdaBoost always converges to a finite cycle. The construction is based on a block-product gadget whose two factors share an exact period-2 orbit for their 5-step branch maps, but whose linearized return maps have dominant eigenvalues with an irrational logarithmic ratio. This irrationality forces the burst-winner sequence to have an irrational asymptotic frequency, precluding eventual periodicity. All assertions are certified by exact rational arithmetic with no floating-point computation in the proof chain. This work was developed in collaboration with GPT-5.4 Pro and Claude Opus 4.6.

1 Introduction

The AdaBoost algorithm [1] defines a discrete dynamical system on the probability simplex in $\mathbb{R}^m$ [2], where a sequence of distributions D_t is updated based on an $m \times N$ matrix $M \in \{-1, +1\}^{m \times N}$ representing the correlations between N weak hypotheses and m training examples. Under the exhaustive selection of weak hypotheses, the next distribution is determined by the mapping $D_{t+1}(i) = D_t(i)/(1 + r_t M_{ij_t})$, where $j_t = \operatorname{argmax}_j (D_t^\top M)_j$ and $r_t = (D_t^\top M)_{j_t}$. While it is proven that D_t converges to a unique fixed point when the weak learning assumption ($r_t \geq c > 0$) is violated, the behavior remains mathematically elusive when the assumption holds. The iterative map is highly discontinuous due to the argmax operation, leaving the long-term asymptotic behavior—specifically, whether the sequence $\{D_t\}$ always enters a cycle of length l such that $D_{kl+b} \to \hat{D}_b$ as $k \to \infty$—an open question [2, 3].

The problem under consideration is the following: prove or disprove that for every $\{-1, +1\}$-valued matrix M, the sequence of distributions D_t computed by the AdaBoost mapping necessarily converges to a cycle. In this paper, we disprove the conjecture by constructing an explicit counterexample.

1.1 Proof outline

We construct a 20×8 sign matrix $M_0 = L_A \boxplus L_B$ as a block product of two small gadgets whose dynamics decouple on an invariant product manifold (Section 3). Both gadgets share an exact period-2 orbit for the 5-step branch maps (Section 6). The linearized return maps at this orbit have dominant eigenvalues λ and κ whose ratio $\log \lambda / \log \kappa$ is irrational, as shown by a norm argument in algebraic number theory (Section 7). Rigorous interval arithmetic certifies explicit invariant sectors around the orbit, together with contraction rates and burst-locking margins (Sections 8–10). Within these sectors, the irrational eigenvalue ratio forces the burst-winner sequence to have an irrational asymptotic frequency, which precludes eventual periodicity and hence finite-cycle convergence (Sections 11–13). A row-duplication argument converts an explicit rational starting point into a uniform-start counterexample on a larger matrix $\widetilde{M}$ with weak-learning margin 1/5 (Section 14).

2 Main result

We work with exhaustive AdaBoost on a fixed sign matrix $M \in \{-1,+1\}^{m\times N}$, with smallest-index tie-breaking. If

$$\mu_t(j) = (D_t^\top M)_j, \qquad j_t = \min \operatorname*{argmax}_{1\le j\le N} \mu_t(j), \qquad r_t = \mu_t(j_t), \tag{2.1}$$

then the update is

$$D_{t+1}(i) = \frac{D_t(i)}{1 + r_t M_{ij_t}}. \tag{2.2}$$

The accompanying certificate uses only exact symbolic algebra over $\mathbf{Q}$, exact interval arithmetic with `fractions.Fraction` endpoints, and exact square-root enclosures obtained by integer arithmetic, so every assertion in the certificate is rigorous. Decimal intervals are used only for reporting.

Theorem 1. *There exists a finite $\{-1,+1\}$-valued matrix $\widetilde{M}$ with weak-learning margin*

$$\gamma(\widetilde{M}) = \frac{1}{5} > 0 \tag{2.3}$$

*such that standard exhaustive AdaBoost on $\widetilde{M}$, started from the exact uniform distribution and using smallest-index tie-breaking, does **not** converge to any finite cycle.*

*Consequently, the COLT question "does exhaustive AdaBoost always converge to a finite cycle?" has a **negative** answer.*

Corollary 2. *There exists a finite $\{-1,+1\}$-valued matrix $\widetilde{M}'$ whose column set is closed under negation (i.e., for every column h of $\widetilde{M}'$, $-h$ also appears as a column of $\widetilde{M}'$), with weak-learning margin*

$$\gamma(\widetilde{M}') = \frac{1}{5} > 0, \tag{2.4}$$

*such that standard exhaustive AdaBoost on $\widetilde{M}'$, started from the exact uniform distribution and using smallest-index tie-breaking, does **not** converge to any finite cycle. In particular, the COLT question in its original hypothesis-class-closed-under-negation formulation also has a negative answer.*

The proof of the Corollary appears at the end of Section 14.

The remainder of the paper is devoted to the proof of Theorem 1, which occupies Sections 3–14.

3 The gadgets, the product system, and the exact invariant affine manifolds

The counterexample is built from two small sign matrices, which we call gadgets. Let

$$L_A = \begin{pmatrix} 1 & 1 & -1 & -1 \\ -1 & 1 & 1 & 1 \\ 1 & -1 & -1 & 1 \\ 1 & -1 & 1 & -1 \end{pmatrix}, \qquad L_B = \begin{pmatrix} 1 & 1 & -1 & -1 \\ -1 & 1 & 1 & 1 \\ 1 & -1 & -1 & 1 \\ 1 & -1 & 1 & -1 \\ 1 & -1 & 1 & 1 \end{pmatrix}. \tag{3.1}$$

Define the 20×8 block-product matrix $M_0 = L_A \boxplus L_B$ as follows. Its rows are indexed by pairs

$$(r, s) \in \{1, 2, 3, 4\} \times \{1, 2, 3, 4, 5\},$$

and its columns are

$$A_1, \ldots, A_4, B_1, \ldots, B_4.$$

The entries are

$$M_0\big((r, s), A_j\big) = (L_A)_{rj}, \qquad M_0\big((r, s), B_j\big) = (L_B)_{sj}.$$

Lemma 3 (reduction to row marginals)**.** *Let $W = (w_{rs})$ be **any** distribution on the 20 rows of M_0, and define its row marginals*

$$a_r = \sum_{s=1}^{5} w_{rs}, \qquad b_s = \sum_{r=1}^{4} w_{rs}. \tag{3.2}$$

Then the edge vector on the A-block depends only on a, and the edge vector on the B-block depends only on b:

$$(W^\top M_0)_{A_j} = (a^\top L_A)_j \qquad (1 \le j \le 4), \tag{3.3}$$

$$(W^\top M_0)_{B_j} = (b^\top L_B)_j \qquad (1 \le j \le 4). \tag{3.4}$$

Proof. For $1 \le j \le 4$,

$$(W^\top M_0)_{A_j} = \sum_{r=1}^{4} \sum_{s=1}^{5} w_{rs} (L_A)_{rj} = \sum_{r=1}^{4} a_r (L_A)_{rj} = (a^\top L_A)_j.$$

The B-block identity is identical. □

Lemma 4 (exact invariance of the product manifold)**.** *Let*

$$\mathcal{P} = \{a \otimes b : \ a \in \Delta^3, \ b \in \Delta^4\}.$$

If $W = a \otimes b \in \mathcal{P}$ and AdaBoost selects an A-column A_j, then the updated state is

$$W' = a' \otimes b,$$

where a' is the one-step exhaustive-AdaBoost update of a on L_A using column j. If AdaBoost selects a B-column B_j, then

$$W' = a \otimes b',$$

where b' is the one-step update of b on L_B using column j.

Hence $\mathcal{P}$ is exactly invariant.

Proof. Assume an A-column j is selected. By Lemma 3, the selected edge is

$$\mu = (a^\top L_A)_j.$$

Therefore

$$w'_{rs} = \frac{a_r b_s}{1 + \mu (L_A)_{rj}} = \left(\frac{a_r}{1 + \mu (L_A)_{rj}} \right) b_s = a'_r b_s,$$

where a' is exactly the one-step AdaBoost update of a on L_A. The B-column case is identical. □

Each 5-step burst follows one of two column-selection sequences, determined by the current phase. The two 5-step half-words are

$$H_0 = (1, 3, 4, 1, 2), \qquad H_1 = (1, 4, 3, 1, 2). \tag{3.5}$$

Lemma 5 (exact affine landing manifolds)**.** *Let w be a distribution on $\{1, \dots, m\}$ and let $h \in \{\pm 1\}^m$ be a selected column such that both sign classes carry positive w-mass:*

$$\sum_{h_i=+1} w_i > 0, \qquad \sum_{h_i=-1} w_i > 0.$$

Equivalently, if $\mu = \sum_i w_i h_i$, then $-1 < \mu < 1$. A single AdaBoost update on h sends the total post-update mass on the $h = +1$ rows and on the $h = -1$ rows to exactly $1/2$ each.

Proof. Write

$$P_+ = \sum_{h_i=+1} w_i, \qquad P_- = \sum_{h_i=-1} w_i, \qquad \mu = \sum_i w_i h_i = P_+ - P_-.$$

Since $P_+ + P_- = 1$, we have $P_+ = (1 + \mu)/2$ and $P_- = (1 - \mu)/2$. Because both sign classes have positive mass, $-1 < \mu < 1$, so the exhaustive AdaBoost update

$$w_i' = \frac{w_i}{1 + \mu h_i}$$

is well-defined. Therefore

$$\sum_{h_i=+1} w_i' = \frac{P_+}{1 + \mu} = \frac{(1 + \mu)/2}{1 + \mu} = \frac{1}{2},$$

and similarly

$$\sum_{h_i=-1} w_i' = \frac{P_-}{1 - \mu} = \frac{(1 - \mu)/2}{1 - \mu} = \frac{1}{2}.$$

□

In all applications below, the current distributions are strictly positive and every relevant column of L_A and L_B contains both signs, so the hypothesis of Lemma 5 is automatic.

Since both H_0 and H_1 end in column 2, every 5-step A-burst lands in

$$\Sigma_A = \{(a, b, c, d) \in \Delta^3 : \ a + b = \tfrac{1}{2}, \ c + d = \tfrac{1}{2}\}, \tag{3.6}$$

and every 5-step B-burst lands in

$$\Sigma_B = \{(a, b, c, d, e) \in \Delta^4 : \ a + b = \tfrac{1}{2}, \ c + d + e = \tfrac{1}{2}\}. \tag{3.7}$$

Write

$$(a, b, c, d) = \Big(p - \tfrac{1}{2}, \ 1 - p, \ \tfrac{1}{2} - d, \ d\Big) \in \Sigma_A, \tag{3.8}$$

$$(a, b, c, d, e) = \Big(p - \tfrac{1}{2}, \ 1 - p, \ \tfrac{1}{2} - d - e, \ d, \ e\Big) \in \Sigma_B. \tag{3.9}$$

The certificate proves exactly that the B third-coordinate numerators are divisible by e, so the face $e = 0$ is invariant and

$$F_0^B|_{e=0} = F_0^A, \qquad F_1^B|_{e=0} = F_1^A. \tag{3.10}$$

Notation convention. Throughout, we distinguish *affine chart points* from *simplex probability vectors* using a hat. We write

$$\hat{P}_i^A = (p_i, d_i) \in \mathbb{R}^2, \qquad \hat{P}_i^B = (p_i, d_i, e_i) \in \mathbb{R}^3$$

for chart coordinates on Σ_A and Σ_B, and

$$P_i^A \in \Delta^3, \qquad P_i^B \in \Delta^4$$

for the corresponding simplex probability vectors obtained from the affine parametrizations above. Similarly, at burst boundary n, we write $\hat{a}_n, \hat{b}_n$ for chart coordinates and a_n, b_n for the simplex vectors, so that $W_n = a_n \otimes b_n$.

4 Margin

We show that the weak-learning margin of the block-product matrix M_0 is exactly $1/5$. The argument uses a single linear witness. Let $\beta = (2, 1, 1, 1)$. Every row r of L_A satisfies $r \cdot \beta = 1$, and every row r of L_B satisfies $r \cdot \beta \in \{1, 3\}$. Thus for every distribution u on either gadget,

$$2\mu_1 + \mu_2 + \mu_3 + \mu_4 = \mathbb{E}_u[r \cdot \beta] \geq 1, \tag{4.1}$$

so $\max_j \mu_j \geq 1/5$. Hence

$$\gamma(L_A), \gamma(L_B) \geq \frac{1}{5}. \tag{4.2}$$

Equality for L_A is attained at

$$u_* = (1/5, 2/5, 1/5, 1/5), \qquad u_*^\top L_A = (1/5, 1/5, 1/5, 1/5). \tag{4.3}$$

It follows that $\gamma(L_A) = 1/5$. Since the first four rows of L_B coincide with L_A, the same minimax argument gives $\gamma(L_B) \leq 1/5$, and together with the lower bound above,

$$\gamma(L_A) = \gamma(L_B) = \frac{1}{5}. \tag{4.4}$$

By Lemma 3, for any distribution W on the rows of M_0 with marginals (a, b),

$$\max_{1 \leq j \leq 8} (W^\top M_0)_j = \max\Big(\max_{1 \leq j \leq 4} (a^\top L_A)_j, \ \max_{1 \leq j \leq 4} (b^\top L_B)_j \Big). \tag{4.5}$$

Conversely, every pair of marginals (a, b) is realized by the product distribution $a \otimes b$. Hence

$$\gamma(M_0) = \min_W \max_j (W^\top M_0)_j = \min_{a,b} \max\Big(\max_j (a^\top L_A)_j, \ \max_j (b^\top L_B)_j \Big) = \frac{1}{5}. \tag{4.6}$$

5 Exact 5-step branch maps

We now compute the closed-form rational maps that result from composing five consecutive AdaBoost updates along each half-word. The two branch maps F_0^G and F_1^G for each gadget $G \in \{A, B\}$ are related by a simple involution. Define the phase-swap maps

$$S_A(p, d) = \big(p, \tfrac{1}{2} - d\big), \qquad S_B(p, d, e) = \big(p, \tfrac{1}{2} - d - e, \ e\big). \tag{5.1}$$

Then the phase-1 branch maps are conjugates of the phase-0 maps:

$$F_1^A = S_A \circ F_0^A \circ S_A, \qquad F_1^B = S_B \circ F_0^B \circ S_B. \tag{5.2}$$

It therefore suffices to give explicit formulas for F_0^A and F_0^B.

A-gadget

The phase-0 branch map on Σ_A is the rational function

$$F_0^A(p,d) = \left(\frac{P_A(p,d)}{Q_A(p,d)}, \frac{D_A(p,d)}{T_A(p,d)} \right), \tag{5.3}$$

where

$$\begin{aligned} P_A = {} & 4d^4 - 6d^3p + 2d^3 - 30d^2p^2 + 15d^2p - 2d^2 \\ & - 6dp^3 + 14dp^2 - 4dp + 6p^4 + p^3 - 2p^2, \end{aligned} \tag{5.4}$$

$$\begin{aligned} Q_A = {} & 8d^4 - 8d^3p + 2d^3 - 48d^2p^2 + 22d^2p - 3d^2 \\ & - 8dp^3 + 22dp^2 - 6dp + 8p^4 + 2p^3 - 3p^2, \end{aligned} \tag{5.5}$$

$$D_A = d\big(2d^3 + 2d^2p - d^2 - 6dp^2 + 2p^3 + p^2\big), \tag{5.6}$$

$$\begin{aligned} T_A = {} & 8d^4 - 2d^3 - 24d^2p^2 + 6d^2p - d^2 \\ & + 10dp^2 - 2dp + 2p^3 - p^2. \end{aligned} \tag{5.7}$$

B-gadget

The phase-0 branch map on Σ_B is

$$F_0^B(p,d,e) = \left(\frac{P_B(p,d,e)}{Q_B(p,d,e)}, \frac{D_B(p,d,e)}{T_B(p,d,e)}, \frac{E_B(p,d,e)}{T_B(p,d,e)} \right), \tag{5.8}$$

where

$$\begin{aligned} P_B = {} & 4d^4 + 16d^3e - 6d^3p + 2d^3 + 20d^2e^2 - 36d^2ep + 12d^2e - 30d^2p^2 + 15d^2p - 2d^2 \\ & + 8de^3 - 46de^2p + 18de^2 - 48dep^2 + 30dep - 4de - 6dp^3 + 14dp^2 - 4dp \\ & - 16e^3p + 8e^3 - 22e^2p^2 + 15e^2p - 2e^2 + 8ep^2 - 4ep + 6p^4 + p^3 - 2p^2, \end{aligned} \tag{5.9}$$

$$\begin{aligned} Q_B = {} & 8d^4 + 32d^3e - 8d^3p + 2d^3 + 40d^2e^2 - 56d^2ep + 16d^2e - 48d^2p^2 + 22d^2p - 3d^2 \\ & + 16de^3 - 72de^2p + 26de^2 - 72dep^2 + 44dep - 6de - 8dp^3 + 22dp^2 - 6dp \\ & - 24e^3p + 12e^3 - 32e^2p^2 + 22e^2p - 3e^2 + 12ep^2 - 6ep + 8p^4 + 2p^3 - 3p^2, \end{aligned} \tag{5.10}$$

$$D_B = d\big(2d^3 + 8d^2e + 2d^2p - d^2 + 10de^2 - 2de - 6dp^2 + 4e^3 - 2e^2p - e^2 - 4ep^2 + 2p^3 + p^2\big), \tag{5.11}$$

$$\begin{aligned} T_B = {} & 8d^4 + 32d^3e - 2d^3 + 40d^2e^2 - 24d^2ep - 24d^2p^2 + 6d^2p - d^2 \\ & + 16de^3 - 32de^2p + 6de^2 - 24dep^2 + 12dep - 2de + 10dp^2 - 2dp \\ & - 8e^3p + 4e^3 - 8e^2p^2 + 6e^2p - e^2 + 4ep^2 - 2ep + 2p^3 - p^2, \end{aligned} \tag{5.12}$$

$$E_B = e\big(2d^3 + 4d^2e - 6d^2p + d^2 + 2de^2 - 8dep + 2de + 2dp^2 - 2e^2p + e^2 + 2p^3 - p^2\big). \tag{5.13}$$

Because E_B factors e exactly, the face $e = 0$ is invariant.

Lemma 6 (regularity of the branch maps)**.** *The denominators Q_A, T_A are nonzero on $\Sigma_A^\circ = \{(p,d) : \frac{1}{2} < p < 1,\ 0 < d < \frac{1}{2}\}$, and Q_B, T_B are nonzero on $\Sigma_B^\circ = \{(p,d,e) : \frac{1}{2} < p < 1,\ d > 0,\ e > 0,\ d+e < \frac{1}{2}\}$. Hence the branch maps F_i^A, F_i^B and all their derivatives are smooth on these domains, and all interval-arithmetic evaluations of the branch maps, their Jacobians, and their Hessians on subsets of these domains involve only nonzero denominators.*

Proof. **Step 1 (sign check).** Every column of L_A contains both $+1$ and -1 entries: column 1 is $(1,-1,1,1)^\top$, column 2 is $(1,1,-1,-1)^\top$, column 3 is $(-1,1,-1,1)^\top$, column 4 is $(-1,1,1,-1)^\top$. The same holds for every column of L_B, since the first four rows of L_B equal L_A and the fifth row adds one further entry per column.

Step 2 (strict positivity implies $|\mu| < 1$). If $a \in \Delta^{m-1}$ is strictly positive and $h \in \{-1,+1\}^m$ has both signs, then

$$\mu = a^\top h = \sum_{h_i=+1} a_i - \sum_{h_i=-1} a_i.$$

Both partial sums are strictly positive, so $-1 < \mu < 1$. Hence every denominator $1 + \mu h_i \in \{1+\mu,\ 1-\mu\}$ is strictly positive, and the updated vector $a_i' = a_i/(1+\mu h_i)$ is again strictly positive with $\sum_i a_i' = 1$.

Step 3 (well-definedness of the five-step composition). A point $(p,d) \in \Sigma_A^\circ$ corresponds to the strictly positive distribution $(p - \frac{1}{2},\ 1-p,\ \frac{1}{2} - d,\ d) \in \Delta^3$. By Steps 1 and 2, each of the five individual AdaBoost updates in the column sequence H_0 (or H_1) has strictly positive denominators and produces a strictly positive output distribution. Hence the five-step composition is well-defined on all of Σ_A°. The identical argument applies to L_B on Σ_B°.

Step 4 (reduced denominators are nonzero). Before algebraic simplification, each coordinate of the five-step composition is a rational function $\tilde{P}/\tilde{Q}$ whose denominator $\tilde{Q}$ is a product of the individual-step denominators (polynomials in the chart coordinates obtained by successive substitution and clearing of fractions). By Step 3, each factor is strictly positive on Σ_A°, so $\tilde{Q} > 0$ there. The reduced-form denominator Q satisfies $Q \mid \tilde{Q}$: since $P\tilde{Q} = Q\tilde{P}$ and $\gcd(P,Q) = 1$, Euclid's lemma in the polynomial ring $\mathbf{Q}[p,d]$ gives $Q \mid \tilde{Q}$. Hence any real zero of Q would also be a zero of $\tilde{Q}$. Since $\tilde{Q}$ has no zeros on Σ_A°, neither does Q. The same argument applies to T_A, Q_B, and T_B.

Step 5 (smoothness and derivatives). Since F_i^A and F_i^B are rational functions with nonzero denominators, they are C^∞ on Σ_A° and Σ_B° respectively. Their Jacobians and Hessians are rational functions whose denominators are powers of Q_A, T_A (resp. Q_B, T_B), hence also nonzero. The phase-1 maps inherit the same regularity, since $S_A : \Sigma_A^\circ \to \Sigma_A^\circ$ and $S_B : \Sigma_B^\circ \to \Sigma_B^\circ$. □

6 Exact common 2-cycle

Both gadgets share an exact period-2 orbit for the 5-step branch maps. A point $\hat{P}_0^A = (p,d)$ lies on a 2-cycle if and only if F_0^A maps it to the phase-swapped point $(p, \frac{1}{2} - d)$, which F_1^A then maps back. Equivalently, the 2-cycle is determined by the fixed-point equations

$$F_0^A(p,d) = \big(p, \tfrac{1}{2} - d\big). \tag{6.1}$$

By Lemma 6, the denominators of F_0^A are nonzero on Σ_A°, so clearing denominators and cancelling factors that are nonzero on Σ_A° is legitimate. The two coordinate equations are equivalent on Σ_A° to

$$\begin{aligned} N_1(p,d) = {} & 2d^4 - 2d^3p + d^3 - 12d^2p^2 + 7d^2p - d^2 \\ & - 2dp^3 + 6dp^2 - 2dp + 2p^4 - p^2, \end{aligned} \tag{6.2}$$

$$
\begin{aligned}
N_2(p,d) = & -16d^5 + 8d^4 + 48d^3p^2 - 16d^3p + 2d^3 \\
& - 32d^2p^2 + 10d^2p - d^2 - 8dp^3 + 10dp^2 - 2dp + 2p^3 - p^2.
\end{aligned} \tag{6.3}
$$

The resultant of N_1 and N_2 with respect to d is

$$
\mathrm{Res}_d(N_1, N_2) = 4096\, p^{10}(2p-1)^3 f(p), \tag{6.4}
$$

where

$$
f(p) = 1024p^7 - 3072p^6 + 3200p^5 - 1152p^4 - 144p^3 + 176p^2 - 40p + 7. \tag{6.5}
$$

The factors $p^{10}(2p-1)^3$ are nonzero on Σ_A° (since $p > \frac{1}{2}$), so any interior solution must satisfy $f(p) = 0$.

The certificate isolates a unique root

$$
p_* \in [0.637116837818846,\ 0.637116837818848], \tag{6.6}
$$

and proves

$$
\begin{aligned}
d_* &= 512p_*^6 - 1152p_*^5 + 736p_*^4 - 24p_*^3 - 88p_*^2 + 19p_* - \tfrac{9}{2} \\
&\in [0.257995382953758869,\ 0.257995382958145447].
\end{aligned} \tag{6.7}
$$

In chart coordinates, set

$$
\hat{P}_0^A = (p_*,\, d_*), \qquad \hat{P}_1^A = (p_*,\, \tfrac{1}{2} - d_*). \tag{6.8}
$$

The corresponding simplex vectors are

$$
P_0^A = \Big(p_* - \tfrac{1}{2},\ 1 - p_*,\ \tfrac{1}{2} - d_*,\ d_*\Big), \qquad P_1^A = \Big(p_* - \tfrac{1}{2},\ 1 - p_*,\ d_*,\ \tfrac{1}{2} - d_*\Big). \tag{6.9}
$$

Then

$$
F_0^A(\hat{P}_0^A) = \hat{P}_1^A, \qquad F_1^A(\hat{P}_1^A) = \hat{P}_0^A. \tag{6.10}
$$

In chart coordinates, set

$$
\hat{P}_0^B = (p_*,\, d_*,\, 0), \qquad \hat{P}_1^B = (p_*,\, \tfrac{1}{2} - d_*,\, 0), \tag{6.11}
$$

with corresponding simplex vectors

$$
P_0^B = (P_0^A, 0), \qquad P_1^B = (P_1^A, 0). \tag{6.12}
$$

Since $F_i^B|_{e=0} = F_i^A$, the same identities hold for B:

$$
F_0^B(\hat{P}_0^B) = \hat{P}_1^B, \qquad F_1^B(\hat{P}_1^B) = \hat{P}_0^B. \tag{6.13}
$$

Thus both gadgets share an exact period-2 orbit for the 5-step branch maps, equivalently a period-10 orbit in AdaBoost time.

7 Exact spectra and the irrationality input

The key algebraic input is that the dominant eigenvalues of the two gadgets' return maps have an irrational logarithmic ratio, which will ultimately force the burst-winner sequence to be aperiodic.

Denote by $J_i^A = DF_i^A(\hat{P}_i^A)$ and $J_i^B = DF_i^B(\hat{P}_i^B)$ the Jacobians of the branch maps at the 2-cycle points, and define the full-period return maps

$$R_A = J_1^A J_0^A, \qquad R_B = J_1^B J_0^B. \tag{7.1}$$

The return map R_B has a block-triangular structure:

$$R_B = \begin{pmatrix} R_A & u \\ 0 & \kappa \end{pmatrix}. \tag{7.2}$$

This block-triangular structure holds because the e-component E_B of each branch map factors as $e \cdot g(p, d, e)$, so at $e = 0$ the partial derivatives $\partial E_B/\partial p$ and $\partial E_B/\partial d$ vanish. Each $J_i^B|_{e=0}$ therefore has zeros in positions $(3, 1)$ and $(3, 2)$, and matrix multiplication preserves this pattern. The upper-left 2×2 block equals R_A because $F_i^B|_{e=0} = F_i^A$.

Define λ as the unique root of the irreducible degree-14 polynomial

$$\begin{aligned} P_\lambda(x) = {} & 16x^{14} - 117232x^{13} - 8330040x^{12} - 3628852593x^{11} \\ & + 766183444916x^{10} - 37440753370248x^9 \\ & + 504085851461616x^8 + 57085439654130x^7 \\ & + 18299281845941356x^6 - 5676298472505780x^5 \\ & + 422743039128068x^4 - 4789675432569x^3 \\ & + 7468158864x^2 - 2899292x + 4 \end{aligned} \tag{7.3}$$

in the isolating interval $[0.0944920594891271,\ 0.0944920594891272]$. Hence

$$N_{\mathbf{Q}(\lambda)/\mathbf{Q}}(\lambda) = \frac{4}{16} = \frac{1}{4}. \tag{7.4}$$

The second eigenvalue of R_A is also a root of P_λ; define ν as the unique root of P_λ in $[0.0011453882842312,\ 0.0011453882842313]$.

Define κ as the unique root of the irreducible degree-7 polynomial

$$P_\kappa(x) = x^7 - 35x^6 + 385x^5 - 1835x^4 + 2315x^3 - 625x^2 + 51x - 1 \tag{7.5}$$

in the isolating interval $[0.2038817328642625,\ 0.2038817328642626]$. Hence

$$N_{\mathbf{Q}(\kappa)/\mathbf{Q}}(\kappa) = 1. \tag{7.6}$$

The isolating intervals give

$$0 < \nu < \lambda < \kappa < 1.$$

If $\log \lambda / \log \kappa \in \mathbf{Q}$, then $\lambda^m = \kappa^n$ for some $m, n > 0$. Taking norms in a common compositum field gives

$$(1/4)^{m[L:\mathbf{Q}(\lambda)]} = 1^{n[L:\mathbf{Q}(\kappa)]}, \tag{7.7}$$

which is impossible. Therefore

$$\frac{\log \lambda}{\log \kappa} \notin \mathbf{Q}. \tag{7.8}$$

8 Exact dominant half-step directions and half-step contractions

The return map $R_G = J_1^G J_0^G$ governs the full-period linearized dynamics, but the proof requires a finer decomposition into half-steps. The phase symmetries from Section 5 relate the two Jacobians at corresponding phases of the 2-cycle. Recall that $F_1^G = S_G \circ F_0^G \circ S_G$ for each gadget. Differentiating at the 2-cycle points yields

$$J_1^A = DS_A\, J_0^A\, DS_A, \qquad J_1^B = DS_B\, J_0^B\, DS_B,$$

where the differentials of the phase-swap maps are the constant matrices

$$DS_A = \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}, \qquad DS_B = \begin{pmatrix} 1 & 0 & 0 \\ 0 & -1 & -1 \\ 0 & 0 & 1 \end{pmatrix}.$$

Lemma 7 (half-step factorization)**.** *For each gadget $G \in \{A, B\}$, let $R_G = J_1^G J_0^G$ be the full return map with dominant eigenvalue μ_G ($\mu_A = \lambda$, $\mu_B = \kappa$). Let v_0^G be a corresponding eigenvector normalized so that its p-component is 1, and set $v_1^G = DS_G\, v_0^G$. Then there exists a unique positive scalar c_G such that*

$$J_0^G v_0^G = c_G v_1^G, \qquad J_1^G v_1^G = c_G v_0^G, \qquad c_G^2 = \mu_G. \tag{8.1}$$

Proof. We give the argument for gadget A; gadget B is identical with DS_B, κ replacing DS_A, λ, once we verify that κ is a simple eigenvalue of R_B.

Step 1: conjugate return map. Define the reverse return map $R'_A = J_0^A J_1^A$. Using $J_1^A = DS_A J_0^A DS_A$ and $DS_A^2 = I$,

$$DS_A R_A DS_A = DS_A J_1^A J_0^A DS_A = (DS_A J_1^A DS_A)(DS_A J_0^A DS_A) = J_0^A J_1^A = R'_A.$$

Hence R'_A is conjugate to R_A and has the same eigenvalues λ, ν. Moreover,

$$R'_A(DS_A v_0^A) = DS_A R_A v_0^A = \lambda\, DS_A v_0^A = \lambda v_1^A,$$

so v_1^A is the λ-eigenvector of R'_A.

Step 2: $w = J_0^A v_0^A$ is a λ-eigenvector of R'_A. From $R_A v_0^A = \lambda v_0^A$ we get $J_1^A w = \lambda v_0^A$. Therefore

$$R'_A w = J_0^A(J_1^A w) = J_0^A(\lambda v_0^A) = \lambda J_0^A v_0^A = \lambda w.$$

Step 3: simplicity forces $w \in \operatorname{span}(v_1^A)$. Since $\lambda \neq \nu$, the λ-eigenspace of R'_A is one-dimensional. Because both w and v_1^A are λ-eigenvectors, there exists a scalar c_A such that

$$J_0^A v_0^A = c_A v_1^A.$$

Step 4: the second half-step identity and $c_A^2 = \lambda$. Using $J_1^A = DS_A J_0^A DS_A$, $v_1^A = DS_A v_0^A$, and $DS_A^2 = I$:

$$J_1^A v_1^A = DS_A J_0^A DS_A (DS_A v_0^A) = DS_A J_0^A v_0^A = DS_A(c_A v_1^A) = c_A v_0^A.$$

Composing the two half-steps gives

$$R_A v_0^A = J_1^A(c_A v_1^A) = c_A(c_A v_0^A) = c_A^2 v_0^A = \lambda v_0^A,$$

so $c_A^2 = \lambda$.

Step 5: positivity from certified half-step p-multipliers. Because the p-component of v_1^A is exactly 1, the scalar c_A is the p-component of $J_0^A v_0^A$. The certificate computes this p-component directly and proves

$$(J_0^A v_0^A)_p,\ (J_1^A v_1^A)_p \in [0.307395607438829821,\ 0.307395607492351897] \subset (0,1).$$

Hence $c_A > 0$. Since $c_A^2 = \lambda$, it follows that $c_A = \sqrt{\lambda}$.

Gadget B: simplicity of κ and positivity. The block-triangular structure of R_B gives eigenvalues λ, ν, κ. By the certified intervals, $\kappa \neq \lambda$ and $\kappa \neq \nu$; hence κ is a simple eigenvalue of the block-triangular R_B, and the same argument applies. The certificate certifies

$$(J_0^B v_0^B)_p \in [0.451532647302682166,\ 0.451532650145556158],$$

$$(J_1^B v_1^B)_p \in [0.451532647288647226,\ 0.451532650159591153] \subset (0,1),$$

so $c_B > 0$, and therefore $c_B = \sqrt{\kappa}$. □

A-gadget eigenvectors and contraction

We now give explicit eigenvector data for each gadget and verify the half-step contraction. Since $\lambda \neq \nu$, the eigenvalue λ of R_A is simple. Define θ_A as the unique scalar such that $R_A(1,\theta_A)^\top = \lambda(1,\theta_A)^\top$, and set $v_0^A = (1,\theta_A)$. The certificate verifies

$$\theta_A \in [0.580405272518281104,\ 0.580405291461442915]. \tag{8.2}$$

Set

$$v_1^A = DS_A v_0^A = (1, -\theta_A). \tag{8.3}$$

By Lemma 7 there is an exact $c_A > 0$ such that

$$J_0^A v_0^A = c_A v_1^A, \qquad J_1^A v_1^A = c_A v_0^A, \qquad c_A^2 = \lambda. \tag{8.4}$$

Hence

$$c_A \in [0.307395607465570497,\ 0.307395607465570659] < 1. \tag{8.5}$$

B-gadget eigenvectors and contraction

Since $\kappa \notin \{\lambda, \nu\}$, the eigenvalue κ of R_B is simple. Define $(\theta_{B,0}, \chi_B)$ as the unique pair such that $R_B(1,\theta_{B,0},\chi_B)^\top = \kappa(1,\theta_{B,0},\chi_B)^\top$, and set $v_0^B = (1,\theta_{B,0},\chi_B)$. The certificate verifies

$$\theta_{B,0} \in [0.198318090249174405,\ 0.198318111627417224], \tag{8.6}$$

$$\chi_B \in [1.01349886578531906,\ 1.01349888527878102]. \tag{8.7}$$

Set

$$\begin{aligned} v_1^B &= DS_B v_0^B = (1, \theta_{B,1}, \chi_B), \\ \theta_{B,1} &= -(\theta_{B,0} + \chi_B) \in [-1.21181699690619825,\ -1.21181695603449347]. \end{aligned} \tag{8.8}$$

By Lemma 7 there is an exact $c_B > 0$ such that

$$J_0^B v_0^B = c_B v_1^B, \qquad J_1^B v_1^B = c_B v_0^B, \qquad c_B^2 = \kappa. \tag{8.9}$$

Hence

$$c_B \in [0.451532648724610500,\ 0.451532648724610611] < 1. \tag{8.10}$$

9 Explicit phase sectors and their certified invariance

We define explicit cone-like sectors around each phase of the 2-cycle and certify their forward invariance under the branch maps via rigorous interval arithmetic.

The sector parameters are

$$r = 2.5 \times 10^{-8}, \qquad \rho = 10^{-3},$$

where r bounds the principal coordinate (distance from the 2-cycle along the dominant eigendirection) and ρ bounds the transverse slope.

A-sectors

For each phase $i \in \{0, 1\}$, the A-sector is

$$S_i^A = \Big\{ \hat{P}_i^A + x\, v_i^A + y\,(0,1) : 0 < x \le r,\ |y| \le \rho x \Big\}. \tag{9.1}$$

B-sectors

The B-sectors have two transverse directions:

$$S_i^B = \Big\{ \hat{P}_i^B + x\, v_i^B + y\,(0,1,0) + z\,(0,0,1) : 0 < x \le r,\ |y| \le \rho x,\ |z| \le \rho x \Big\}. \tag{9.2}$$

Because $\chi_B - \rho > 1$, every point of S_i^B has $e > 0$, so $S_i^B \subset \Sigma_B^\circ$.

Lemma 8 (exact first-column structure in the adapted bases)**.** *Define the adapted basis matrices*

$$B_0^A = [v_0^A, (0,1)], \qquad B_1^A = [v_1^A, (0,1)], \tag{9.3}$$

$$B_0^B = [v_0^B, (0,1,0), (0,0,1)], \qquad B_1^B = [v_1^B, (0,1,0), (0,0,1)]. \tag{9.4}$$

Then, for $i \in \{0, 1\}$,

$$(B_{i\oplus 1}^A)^{-1} J_i^A B_i^A = \begin{pmatrix} c_A & * \\ 0 & * \end{pmatrix}, \tag{9.5}$$

and

$$(B_{i\oplus 1}^B)^{-1} J_i^B B_i^B = \begin{pmatrix} c_B & * & * \\ 0 & * & * \\ 0 & * & * \end{pmatrix}. \tag{9.6}$$

Proof. The first column of B_i^A is v_i^A. By Lemma 7,

$$J_i^A v_i^A = c_A v_{i\oplus 1}^A.$$

Expressing this in the codomain basis $B_{i\oplus 1}^A$ gives first coordinate c_A and zero transverse coordinate. That is exactly the displayed matrix form. The B-case is identical, using $J_i^B v_i^B = c_B v_{i\oplus 1}^B$. □

The certificate computes interval Jacobians in these exact bases and exact second-derivative bounds on explicit rational phase boxes covering the sectors. The B-phase boxes extend to $e = 0$ (their e-coordinate ranges over $[0, k_{Be} r]$); since $Q_B|_{e=0} = Q_A$ and $T_B|_{e=0} = T_A$, Lemma 6 (applied to the A-gadget at $e = 0$ and to the B-gadget for $e > 0$) ensures that all denominators are nonzero on the full closed phase boxes, so every interval-arithmetic division is well-defined. By Lemma 8, the lower entries in the first transformed column vanish **exactly**, so the sector inequalities only use the remaining coefficients and the certified Hessian remainder terms. The resulting certified half-step bounds are:

A-sector certified bounds

In the sector coordinates, the ratio x'/x measures the contraction factor along the dominant eigendirection (and should be close to the half-step eigenvalue $c_A \approx 0.3074$), while $|y'|/x'$ measures the transverse drift relative to the cone boundary ρ.

For both phases $i = 0, 1$,

$$\frac{x'}{x} \in [0.307394561088771942,\ 0.307396653842409774], \tag{9.7}$$

$$\frac{|y'|}{x'} \leq 1.10321761561807761 \times 10^{-4} < \rho. \tag{9.8}$$

(For phase 0 the certified bound is slightly smaller, $1.10321759830186699 \times 10^{-4}$.)

B-sector certified bounds

At phase 0:

$$\frac{x'}{x} \in [0.451388989838190632,\ 0.451676307610047724], \tag{9.9}$$

$$\frac{\max(|y'|, |z'|)}{x'} \leq 6.82589234315899433 \times 10^{-4} < \rho. \tag{9.10}$$

At phase 1:

$$\frac{x'}{x} \in [0.451389955166172293,\ 0.451675342282066063], \tag{9.11}$$

$$\frac{\max(|y'|, |z'|)}{x'} \leq 6.84951234683088803 \times 10^{-4} < \rho. \tag{9.12}$$

The A-gadget bounds above are uniform over both phases (the phase-0 bound on $|y'|/x'$ is slightly tighter, as noted in the parenthetical), so they are reported together. The B-gadget's 3-dimensional dynamics make the two phases slightly asymmetric, so we report them separately.

Therefore

$$F_i^A(S_i^A) \subset S_{i\oplus 1}^A, \qquad F_i^B(S_i^B) \subset S_{i\oplus 1}^B. \tag{9.13}$$

The same computation yields explicit transverse-ratio recurrences, which will be needed for the asymptotic analysis in Section 12:

$$\tau_{m+1}^A \leq \alpha_A \tau_m^A + \beta_A x_m^A, \tag{9.14}$$

$$\alpha_A \leq 0.110098192661922619, \quad \beta_A \leq 8.94275599540566418,$$

$$\tau_{m+1}^B \leq \alpha_B \tau_m^B + \beta_B x_m^B, \tag{9.15}$$

$$\alpha_B \leq 0.684664374851195431, \quad \beta_B \leq 12.8511069222196774,$$

where $\tau_m^A = |y_m|/x_m^A$ and $\tau_m^B = \max(|y_m|, |z_m|)/x_m^B$. Since $\alpha_A, \alpha_B < 1$, the transverse ratios are summably small along the local update subsequences.

10 Local branch words and whole-burst locking

Within the invariant sectors, we must verify two properties: that the column-selection sequence within each gadget matches the expected half-word (H_0 or H_1), and that once a gadget wins the first round of a 5-step burst, it wins all five rounds. Both are established by certified gap estimates.

The certificate checks the selected edge minus the best competing local edge on the entire phase boxes covering the sectors. The minimum certified local branch gap is

$$g_{\mathrm{br}}^{\mathrm{box}} > 0.0250983508506417014. \tag{10.1}$$

It also certifies that at steps $2, 3, 4, 5$ of each half-word the selected edge exceeds $\mu_* = 2p_* - 1$ by at least

$$\Delta_{\mathrm{box}} > 0.0758691457701749199. \tag{10.2}$$

The largest possible first-edge surplus on any phase box is only

$$\delta_{\mathrm{first}}^{\max} < 1.05687287632489371 \times 10^{-7}, \tag{10.3}$$

and consequently the whole-burst lock margin satisfies

$$\Delta_{\mathrm{box}} - \delta_{\mathrm{first}}^{\max} > 0.0758690400828872874. \tag{10.4}$$

These estimates have three immediate consequences. First, the local A-word is exactly H_0 on S_0^A and exactly H_1 on S_1^A; likewise, the local B-word is exactly H_0 on S_0^B and exactly H_1 on S_1^B. Second, and more importantly, once a gadget wins the first round of a burst, it stays globally ahead for the remaining four rounds of that burst. To see this, observe that during an A-burst the B-state does not move at all, so the best B-correlation throughout the burst is bounded by its first-round best edge; the same holds with A and B reversed. Since the strict margin

$$\Delta_{\mathrm{box}} - \delta_{\mathrm{first}}^{\max} > 0 \tag{10.5}$$

separates the selected edge from the first-edge surplus, the first-round winner necessarily keeps winning for rounds $2, 3, 4, 5$.

11 Product burst dynamics

Because the starting point lies in $S_0^A \times S_0^B$ (Section 14), the certified sector invariance $F_i^G(S_i^G) \subset S_{i\oplus 1}^G$ from Section 9 and the whole-burst lock from Section 10 imply by induction that $\hat{a}_n \in S_{i_n}^A$ and $\hat{b}_n \in S_{j_n}^B$ for every burst boundary n.

Let

$$W_n = a_n \otimes b_n \in \mathcal{P} \tag{11.1}$$

be the burst-boundary states, where $a_n \in \Delta^3$ and $b_n \in \Delta^4$ are the simplex marginals. Write the corresponding chart coordinates as

$$\hat{a}_n = \hat{P}_{i_n}^A + x_n^A v_{i_n}^A + y_n^A(0,1), \tag{11.2}$$

$$\hat{b}_n = \hat{P}_{j_n}^B + x_n^B v_{j_n}^B + y_n^B(0,1,0) + z_n^B(0,0,1), \tag{11.3}$$

with $\hat{a}_n \in S_{i_n}^A$ and $\hat{b}_n \in S_{j_n}^B$. For brevity write

$$x_n := x_n^A, \qquad s_n := x_n^B.$$

At the start of a burst, the first selected local edge in either gadget is exactly

$$2p - 1 = \mu_* + 2(p - p_*).$$

Since in the sector coordinates $p - p_* = x$, the first global round chooses

- A if $x_n > s_n$;
- B if $s_n > x_n$;
- A on equality, by smallest-index tie-breaking, because the A-columns have smaller global indices than the B-columns.

Because of the whole-burst lock margin from Section 10, that first-round winner then wins the entire 5-step burst.

Half-step contraction bounds

The certified sector bounds from Section 9 yield the following contraction estimates. If A wins burst n, then

$$x_{n+1} \in [m_A, M_A]\, x_n, \qquad s_{n+1} = s_n, \tag{11.4}$$

with

$$m_A = 0.307394561088771942, \qquad M_A = 0.307396653842409774. \tag{11.5}$$

If B wins burst n, then

$$s_{n+1} \in [m_B, M_B]\, s_n, \qquad x_{n+1} = x_n, \tag{11.6}$$

with

$$m_B = 0.451388989838190632, \qquad M_B = 0.451676307610047724. \tag{11.7}$$

Writing $r_n = x_n/s_n$ for the ratio of the two principal coordinates, the contraction bounds above translate to: if A wins burst n, then $r_{n+1} \in [m_A r_n,\, M_A r_n]$; if B wins, then $r_{n+1} \in [r_n/M_B,\, r_n/m_B]$.

Since A wins precisely when $r_n \geq 1$ and B wins when $r_n < 1$, the ratio r_n undergoes a simple multiplicative random walk whose dynamics we now summarize.

Lemma 9 (burst-ratio dynamics and run-length bounds)**.** *For all sufficiently large n, one has $r_n \in [m_A,\, 1/m_B]$. Moreover:*

(i) every A-run has length exactly 1;

(ii) every B-run has length at most 2.

Consequently, every block of three consecutive bursts contains at least one A-burst and at least one B-burst, and

$$x_n + s_n \to 0 \tag{11.8}$$

exponentially.

Proof. **Step 1 (entry into the compact interval).** After a B-burst at step n (so $r_n < 1$), we have $r_{n+1} \leq r_n/m_B < 1/m_B$. After an A-burst at step n (so $r_n \geq 1$), we have $r_{n+1} \geq m_A r_n \geq m_A$. Once the trajectory has experienced at least one A-burst and one B-burst, r_n enters and remains in $[m_A, 1/m_B]$. For the chosen start (Section 14), $x_A < x_B$, so $r_0 < 1$ and the first burst is B. Each B-burst multiplies s by a factor in $[m_B, M_B] \subset (0, 1)$ while leaving x fixed, so r_n increases geometrically until an A-burst occurs after finitely many steps.

Step 2 (A-runs have length 1). Suppose A wins at step n, so $r_n \geq 1$, and that $r_n \leq 1/m_B$ (since we are in the compact interval). Then

$$r_{n+1} \leq M_A r_n \leq \frac{M_A}{m_B} < 1,$$

where the last inequality is the certified bound $M_A/m_B < 1$. Hence B wins at step $n+1$, and every A-run has length exactly 1.

Step 3 (B-runs have length at most 2). Suppose B wins at step n, so $r_n < 1$, and that $r_n \geq m_A$. After one B-burst,

$$r_{n+1} \geq \frac{r_n}{M_B} \geq \frac{m_A}{M_B}.$$

If $r_{n+1} \geq 1$, the B-run ends after one step. If $r_{n+1} < 1$, then B wins again, and

$$r_{n+2} \geq \frac{r_{n+1}}{M_B} \geq \frac{m_A}{M_B^2} > 1,$$

where the last inequality is the certified bound $m_A/M_B^2 > 1$. Hence A wins at step $n+2$, and every B-run has length at most 2.

Step 4 (exponential convergence). By Steps 2 and 3, every block of three consecutive bursts contains at least one A-burst and at least one B-burst, so

$$x_{n+3} + s_{n+3} \leq \theta\,(x_n + s_n), \qquad \theta = \max(M_A, M_B) < 1.$$

The certified numerical values are

$$\frac{M_A}{m_B} = 0.681001665442930328 < 1,$$
$$\frac{m_A}{M_B^2} = 1.50675124144050170 > 1.$$

□

12 Local asymptotics on the two gadgets

We now establish that the principal coordinates along each gadget's dominant eigendirection contract at rates asymptotically governed by $c_A = \sqrt{\lambda}$ and $c_B = \sqrt{\kappa}$, with summable relative errors. Let $m \mapsto \hat{a}^{(m)}$ be the chart coordinates of the local A-state immediately before the mth A-burst, and $m \mapsto \hat{b}^{(m)}$ the chart coordinates of the local B-state immediately before the mth B-burst. Write

$$x_m^A = p(\hat{a}^{(m)}) - p_*, \qquad x_m^B = p(\hat{b}^{(m)}) - p_*. \tag{12.1}$$

Because the exact basis vectors from Section 8 satisfy

$$J_i^A v_i^A = c_A v_{i\oplus 1}^A, \qquad J_i^B v_i^B = c_B v_{i\oplus 1}^B, \tag{12.2}$$

the certified Taylor bounds from Section 9 give

$$x_{m+1}^A = c_A x_m^A\big(1 + \varepsilon_m^A\big), \qquad |\varepsilon_m^A| \leq \Gamma_A \tau_m^A + \Delta_A x_m^A, \tag{12.3}$$
$$x_{m+1}^B = c_B x_m^B\big(1 + \varepsilon_m^B\big), \qquad |\varepsilon_m^B| \leq \Gamma_B \tau_m^B + \Delta_B x_m^B, \tag{12.4}$$

with certified constants

$$\Gamma_A \le 0.003273115711042833, \qquad \Delta_A \le 5.232171784223787, \tag{12.5}$$
$$\Gamma_B \le 0.31803397447249865, \qquad \Delta_B \le 8.463704731868965, \tag{12.6}$$

and therefore uniform relative-error bounds

$$\Gamma_A \tau + \Delta_A x \le 0.0000034039200056484279 1 \qquad \text{throughout the A sectors,} \tag{12.7}$$
$$\Gamma_B \tau + \Delta_B x \le 0.00031824556709079535 0 \qquad \text{throughout the B sectors.} \tag{12.8}$$

In particular, the relative errors are everywhere < 0.01.

Since $x^A_{m+1} \le M_A x^A_m$ and $x^B_{m+1} \le M_B x^B_m$, both local x-sequences are summable. Iterating the transverse recurrence $\tau_{m+1} \le \alpha\tau_m + \beta x_m$ gives $\tau_m \le \alpha^m \tau_0 + \beta \sum_{k=0}^{m-1} \alpha^{m-1-k} x_k$. Summing over m and exchanging the order of summation,

$$\sum_m \tau_m \le \frac{\tau_0}{1-\alpha} + \frac{\beta}{1-\alpha} \sum_k x_k < \infty. \tag{12.9}$$

In particular, both transverse-ratio series $\sum_m \tau^A_m$ and $\sum_m \tau^B_m$ converge, and consequently

$$\sum_m |\varepsilon^A_m| < \infty, \qquad \sum_m |\varepsilon^B_m| < \infty. \tag{12.10}$$

Since $|\varepsilon^A_m|, |\varepsilon^B_m| < 0.01$ throughout, the bound $|\log(1+u)| \le 2|u|$ applies, giving

$$\log x^A_m = \log x^A_0 + m \log c_A + O(1), \tag{12.11}$$
$$\log x^B_m = \log x^B_0 + m \log c_B + O(1). \tag{12.12}$$

At burst boundary n, the A-state has been updated exactly $N_A(n)$ times and the B-state exactly $N_B(n)$ times, so

$$\log x_n = \log x_0 + N_A(n) \log c_A + O(1), \tag{12.13}$$
$$\log s_n = \log s_0 + N_B(n) \log c_B + O(1). \tag{12.14}$$

Because $r_n = x_n/s_n$ stays in the fixed compact interval $[m_A, 1/m_B]$, the logarithmic difference is bounded:

$$N_A(n) \log c_A - N_B(n) \log c_B = O(1). \tag{12.15}$$

Since $N_A(n) + N_B(n) = n$, it follows that

$$\frac{N_A(n)}{n} \to \frac{-\log c_B}{-\log c_A - \log c_B} = \frac{-\log \kappa}{-\log \lambda - \log \kappa}. \tag{12.16}$$

Because $\log \lambda / \log \kappa \notin \mathbf{Q}$ (Section 7), this limiting frequency is irrational. Since every eventually periodic binary sequence has rational limiting frequency, the burst-winner sequence is **not** eventually periodic.

13 No finite-cycle convergence

We now combine the preceding results to show that the AdaBoost orbit on M_0 does not converge to any finite cycle. The argument proceeds by contradiction: assuming cycle convergence, we deduce that the burst-winner sequence is eventually periodic, contradicting its irrational asymptotic frequency.

Define

$$Q_{ij} = P_i^A \otimes P_j^B, \qquad i, j \in \{0, 1\}. \tag{13.1}$$

These four points are distinct: the tensor product $P_i^A \otimes P_j^B$ determines its marginals by summation, and $P_0^A \neq P_1^A$, $P_0^B \neq P_1^B$ (their 3rd and 4th coordinates are swapped, with $d_* \neq \frac{1}{4}$).

Since $x_n, s_n \to 0$, the burst-boundary states satisfy

$$\|W_n - Q_{i_n j_n}\| \to 0. \tag{13.2}$$

Suppose for contradiction that the full AdaBoost trajectory converges to a cycle of length ℓ, meaning

$$D_{k\ell+b} \to \widehat{D}_b \qquad (k \to \infty) \tag{13.3}$$

for each $b \in \{0, \ldots, \ell - 1\}$. Each burst has exactly 5 steps; since the initial distribution is D_1 (following the COLT convention), the burst-boundary states are $W_n = D_{5n+1}$. Set

$$L = \ell / \gcd(5, \ell). \tag{13.4}$$

Then for each residue $r \in \{0, \ldots, L - 1\}$, the subsequence W_{kL+r} converges to a single limit $\widehat{W}_r$ as $k \to \infty$.

Let

$$\delta = \min\{\|Q - Q'\| : Q, Q' \in \{Q_{00}, Q_{01}, Q_{10}, Q_{11}\},\ Q \neq Q'\} > 0. \tag{13.5}$$

Fix a residue class r, and abbreviate

$$Q_k := Q_{i_{kL+r}, j_{kL+r}}. \tag{13.6}$$

Choose $\varepsilon = \delta/4$. Because $W_{kL+r} \to \widehat{W}_r$ and $\|W_n - Q_{i_n j_n}\| \to 0$, there is K_r such that for all $k \geq K_r$,

$$\|W_{kL+r} - \widehat{W}_r\| < \varepsilon, \qquad \|W_{kL+r} - Q_k\| < \varepsilon. \tag{13.7}$$

Now take any $k, k' \geq K_r$. Then

$$\|Q_k - Q_{k'}\| \leq \|Q_k - W_{kL+r}\| + \|W_{kL+r} - \widehat{W}_r\| + \|\widehat{W}_r - W_{k'L+r}\| + \|W_{k'L+r} - Q_{k'}\| < 4\varepsilon = \delta. \tag{13.8}$$

But distinct points among $\{Q_{00}, Q_{01}, Q_{10}, Q_{11}\}$ are at least δ apart. Therefore $Q_k = Q_{k'}$ for all $k, k' \geq K_r$. So the phase label (i_{kL+r}, j_{kL+r}) is eventually constant on each residue class r.

Hence the phase-label sequence (i_n, j_n) is eventually periodic. Since exactly one coordinate toggles at each burst, the winner of burst n is determined by the transition

$$(i_n, j_n) \longrightarrow (i_{n+1}, j_{n+1}), \tag{13.9}$$

so the burst-winner sequence is also eventually periodic. That contradicts Section 12, which showed that the burst-winner sequence has an irrational limiting frequency and therefore cannot be eventually periodic.

Therefore the AdaBoost orbit on M_0 does **not** converge to any finite cycle.

14 Explicit rational start and explicit uniform-start counterexample

It remains to exhibit a concrete starting point inside the invariant sectors and to convert it into a uniform-start counterexample. Take

$$A = (137116847818847,\ 362883152181153,\ 242004611239991,\ 257995388760009), \tag{14.1}$$

$$B = (137116857552466,\ 362883142447534,\ 242004593130510,\ 257995386869490,\ 20000000). \tag{14.2}$$

Since $\sum_i A_i = \sum_j B_j = 10^{15}$, set

$$a_0 = \frac{A}{10^{15}}, \qquad b_0 = \frac{B}{10^{15}}, \qquad W_0 = a_0 \otimes b_0. \tag{14.3}$$

To apply the sector-based analysis of Sections 9–13, we need to verify that this starting point lies inside the invariant sectors.

The certificate proves that the chart images satisfy

$$\hat{a}_0 \in S_0^A, \qquad \hat{b}_0 \in S_0^B. \tag{14.4}$$

The starting point is decomposed in sector coordinates: we report the principal coordinate x (distance from the 2-cycle along the dominant eigendirection) and the transverse ratios (deviation from the eigendirection, normalized by x), confirming that they satisfy the sector constraint $|\text{transverse ratio}| \le \rho$.

More explicitly:

$$x_A \in [9.999999 \times 10^{-9},\ 1.0000001 \times 10^{-8}],$$

with A-sector residual in

$$[-2.18994134061382760 \times 10^{-4},\ 2.19798676726244218 \times 10^{-4}],$$

and

$$x_B \in [1.9733618 \times 10^{-8},\ 1.9733620 \times 10^{-8}],$$

with B d-residual in

$$[-1.10962941057972692 \times 10^{-4},\ 1.11368115987132873 \times 10^{-4}],$$

and B e-residual in

$$[-9.48895873525338312 \times 10^{-8},\ 2.73218648306331209 \times 10^{-8}],$$

all strictly inside the sector slope bound $\rho = 10^{-3}$. Note that $x_A \approx 10^{-8} < 1.97 \times 10^{-8} \approx x_B$, so $r_0 = x_A/x_B < 1$ and the first burst is a B-burst, as used in Lemma 9's proof.

Lemma 10 (row duplication and class aggregation)**.** *Let M' be obtained from a matrix M by duplicating rows. Partition the rows of M' into duplicate classes, one class for each row of M. Suppose the initial distribution on M' is constant on each duplicate class.*

Then:

1. *the AdaBoost orbit on M' remains constant on each duplicate class for all time;*

2. *the class-summed weights evolve* ***exactly*** *as exhaustive AdaBoost on M started from the corresponding aggregated distribution;*

3. $\gamma(M') = \gamma(M)$.

Proof. Identical rows have the same sign in every column, so at every AdaBoost step they are multiplied by the same factor $1/(1 + r_t h_i)$. Hence equality within each duplicate class is preserved forever.

Let $\bar{D}_t$ be the distribution obtained by summing D'_t over each duplicate class. Because every row in a given class has the same entries in every column, the edge vector of D'_t equals the edge vector of $\bar{D}_t$. Therefore the selected column is the same for D'_t and $\bar{D}_t$, and the class sums satisfy exactly the same update formula as exhaustive AdaBoost on M.

For the margin, every distribution on M' aggregates to a distribution on M with the same edge vector, so $\gamma(M') \geq \gamma(M)$. Conversely, any distribution on M lifts to a class-constant distribution on M' with the same edge vector, so $\gamma(M') \leq \gamma(M)$. Hence $\gamma(M') = \gamma(M)$. □

Now construct $\widetilde{M}$ by replacing the row (i, j) of M_0 by exactly

$$A_i B_j$$

identical copies. Since $\sum_i A_i \sum_j B_j = 10^{30}$, the exact uniform distribution on $\widetilde{M}$ aggregates to W_0. By Lemma 10, the uniform-start AdaBoost orbit on $\widetilde{M}$ projects exactly to the W_0-started orbit on M_0. Since the latter does not converge to a finite cycle, neither does the former.

By Lemma 10 and Section 4,

$$\gamma(\widetilde{M}) = \gamma(M_0) = \frac{1}{5}. \tag{14.5}$$

This completes the proof of Theorem 1. □

Proof of Corollary 2. Set $\widetilde{M}' = [\widetilde{M} \mid -\widetilde{M}]$, where the columns of $-\widetilde{M}$ are appended after those of $\widetilde{M}$. By construction, the column set of $\widetilde{M}'$ is closed under negation.

We verify that no negated column is ever selected, so that the dynamics on $\widetilde{M}'$ are identical to those on $\widetilde{M}$. By Lemma 10, the edge vectors on $\widetilde{M}$ equal those on M_0 at every step, so it suffices to verify on M_0 that the selected column edge strictly exceeds every negated column edge at every step. The certificate does so by exact interval arithmetic on the full sector boxes (items 15–16 of Appendix A). During any burst, two types of negated columns could compete: those from the active gadget and those from the frozen (non-active) gadget. For the frozen gadget, Lemma 3 implies its column edges do not change during the burst; the certificate encloses all frozen-gadget column edges on each sector box. For the active gadget, at each of the five steps within each burst word on each sector box, the certificate encloses all column edges. At every step of every burst/frozen-phase combination, the certificate verifies that the selected edge strictly exceeds every negated column edge. The minimum certified gap over all eight burst/frozen-phase combinations and all five steps is

$$g_{\text{neg}}^{\text{box}} \geq 0.274233539263122735 > 0.$$

Hence no negated column is ever in the argmax set, so the dynamics on $\widetilde{M}'$ are identical to those on $\widetilde{M}$, which by the Theorem does not converge to any finite cycle.

For the margin, the minimax distributions $u_*^A = (1/5, 2/5, 1/5, 1/5)$ on L_A and $u_*^B = (1/5, 2/5, 1/5, 1/5, 0)$ on L_B each give all four column edges equal to exactly $1/5$. Hence every negated column edge is $-1/5$, and $\gamma(\widetilde{M}') = \gamma(\widetilde{M}) = 1/5$. □

A What the certificate proves exactly

The certificate proves all of the following finite assertions.

1. The explicit branch maps $F_0^A, F_1^A, F_0^B, F_1^B$ are derived symbolically from H_0, H_1, and the displayed numerator/denominator polynomials $(P_A, Q_A, D_A, T_A, P_B, Q_B, D_B, T_B, E_B)$ are verified to match the derived rational functions term-for-term.

2. $F_0^B|_{e=0} = F_0^A$, $F_1^B|_{e=0} = F_1^A$, and both B third-coordinate numerators are exactly divisible by e.

3. On the interior region $p > \frac{1}{2}$, the 2-cycle equations are equivalent to $N_1 = N_2 = 0$; the displayed polynomials N_1, N_2 are verified to match the derived expressions term-for-term. The resultant with respect to d is exactly $4096\, p^{10}(2p-1)^3 f(p)$, so the relevant interior factor is the exact degree-7 polynomial f, and the root interval for p_* is unique.

4. The exact polynomial formula for d_* is correct modulo f.

5. The exact elimination polynomial for λ is the irreducible degree-14 polynomial above.

6. The exact elimination polynomial for κ is the irreducible degree-7 polynomial above.

7. The narrow isolating intervals for λ, ν, κ each contain exactly one root.

8. The exact dominant half-step directions v_i^A, v_i^B are certified by explicit rational interval enclosures of $\theta_A, \theta_{B,0}, \chi_B, \theta_{B,1}$.

9. The half-step contractions c_A, c_B are certified and satisfy $c_A < 1$, $c_B < 1$.

10. The explicit sectors S_i^A, S_i^B with $r = 2.5 \times 10^{-8}$ and $\rho = 10^{-3}$ are certified by interval Jacobians and interval Hessian bounds, using the exact first-column structure from Lemma 8.

11. The local branch words H_0, H_1 are certified on the full phase boxes covering those sectors.

12. The whole-burst lock margin is certified.

13. The explicit rational starting point $W_0 = a_0 \otimes b_0$ is certified: its chart image satisfies $\hat{a}_0 \in S_0^A$, $\hat{b}_0 \in S_0^B$.

14. No floating-point number is used in any proof assertion. Floating-point/decimal conversion is used only for human-readable printing of already-certified rational intervals.

15. No negated column is ever selected on $[\widetilde{M} \mid -\widetilde{M}]$: the certificate computes, by exact interval arithmetic on the full sector boxes covering all eight burst/frozen-phase combinations and all five burst steps, interval enclosures of every column edge at every step, and verifies that the selected edge strictly exceeds every negated column edge. The certified minimum gap is $g_{\text{neg}}^{\text{box}} \geq 0.274233539263122735 > 0$.

16. The margin $\gamma([\widetilde{M} \mid -\widetilde{M}]) = 1/5$ is certified: the explicit minimax distributions $u_*^A = (1/5, 2/5, 1/5, 1/5)$ and $u_*^B = (1/5, 2/5, 1/5, 1/5, 0)$ are verified to give all column edges exactly $1/5$ on L_A and L_B respectively, so every negated column edge is $-1/5$.

References

[1] Y. Freund and R. E. Schapire. A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of Computer and System Sciences*, 55(1):119–139, 1997.

[2] C. Rudin, I. Daubechies, and R. E. Schapire. The dynamics of AdaBoost: Cyclic behavior and convergence of margins. *Journal of Machine Learning Research*, 5(Dec):1557–1595, 2004.

[3] C. Rudin, R. E. Schapire, and I. Daubechies. Open problem: Does AdaBoost always cycle? In *Proceedings of the 25th Annual Conference on Learning Theory (COLT)*, volume 23 of *Proceedings of Machine Learning Research*, pages 46.1–46.4. PMLR, 2012.